

# Data Visualization - Online Retail Store

## Forage - Tata Data Visualisation: Empowering Business with Effective Insights

By *Ammaar Ahmad Khan*

### 1. Importing Libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
```

### 2. Loading Dataset

```
df = pd.read_excel('Online_Retail.xlsx', engine='openpyxl')
```

```
print(df)
```

	InvoiceNo	StockCode	Description
Quantity \			
0	536365	85123A	WHITE HANGING HEART T-LIGHT HOLDER
6			
1	536365	71053	WHITE METAL LANTERN
6			
2	536365	84406B	CREAM CUPID HEARTS COAT HANGER
8			
3	536365	84029G	KNITTED UNION FLAG HOT WATER BOTTLE
6			
4	536365	84029E	RED WOOLLY HOTTIE WHITE HEART.
6			
...	...	...	...
...			
541904	581587	22613	PACK OF 20 SPACEBOY NAPKINS
12			
541905	581587	22899	CHILDREN'S APRON DOLLY GIRL
6			
541906	581587	23254	CHILDRENS CUTLERY DOLLY GIRL
4			
541907	581587	23255	CHILDRENS CUTLERY CIRCUS PARADE
4			
541908	581587	22138	BAKING SET 9 PIECE RETROSPOT
3			

	InvoiceDate	UnitPrice	CustomerID	Country
0	2010-12-01 08:26:00	2.55	17850.0	United Kingdom
1	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
2	2010-12-01 08:26:00	2.75	17850.0	United Kingdom
3	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
4	2010-12-01 08:26:00	3.39	17850.0	United Kingdom
...	...	...	...	...
541904	2011-12-09 12:50:00	0.85	12680.0	France
541905	2011-12-09 12:50:00	2.10	12680.0	France
541906	2011-12-09 12:50:00	4.15	12680.0	France
541907	2011-12-09 12:50:00	4.15	12680.0	France
541908	2011-12-09 12:50:00	4.95	12680.0	France

[541909 rows x 8 columns]

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 541909 entries, 0 to 541908
Data columns (total 8 columns):
#   Column          Non-Null Count  Dtype
---  -
0   InvoiceNo        541909 non-null object
1   StockCode       541909 non-null object
2   Description     540455 non-null object
3   Quantity        541909 non-null int64
4   InvoiceDate      541909 non-null datetime64[ns]
5   UnitPrice       541909 non-null float64
6   CustomerID      406829 non-null float64
7   Country         541909 non-null object
dtypes: datetime64[ns](1), float64(2), int64(1), object(4)
memory usage: 33.1+ MB
None
```

```
#General Statistics:
print(df.describe())
```

	Quantity	InvoiceDate	UnitPrice
count	541909.000000	541909	541909.000000
mean	9.552250	2011-07-04 13:34:57.156386048	4.611114
min	-80995.000000	2010-12-01 08:26:00	-11062.060000
25%	1.000000	2011-03-28 11:34:00	1.250000
50%	3.000000	2011-07-19 17:17:00	2.080000
75%	10.000000	2011-10-19 11:27:00	4.130000
max	80995.000000	2011-12-09 12:50:00	38970.000000
std	218.081158	NaN	96.759853

	CustomerID
count	406829.000000
mean	15287.690570

```
min      12346.000000
25%      13953.000000
50%      15152.000000
75%      16791.000000
max      18287.000000
std       1713.600303
```

*#Check for missing values*

```
print("Number of missing values in each column:")
print(df.isnull().sum())
```

Number of missing values in each column:

```
InvoiceNo      0
StockCode      0
Description    1454
Quantity      0
InvoiceDate    0
UnitPrice      0
CustomerID    135080
Country        0
dtype: int64
```

```
print(df.columns)
```

```
Index(['InvoiceNo', 'StockCode', 'Description', 'Quantity',
      'InvoiceDate',
      'UnitPrice', 'CustomerID', 'Country'],
      dtype='object')
```

### 3. Data Cleaning

```
df_raw = df.copy()                                # full dataset
df_clean = df_raw.copy()

# Clean only the copy

# 1. Drop rows with missing Description
df_clean = df_clean.dropna(subset=['Description'])

# 2. Remove cancelled invoices (InvoiceNo starts with 'C')
df_clean =
df_clean[~df_clean['InvoiceNo'].astype(str).str.startswith('C')]

# 3. Remove rows with negative or zero Quantity or UnitPrice
df_clean = df_clean[(df_clean['Quantity'] > 0) &
(df_clean['UnitPrice'] > 0)]

# 4. Add Revenue column
df_clean['Revenue'] = df_clean['Quantity'] * df_clean['UnitPrice']
```

## 4. General Business Metrics

### 4.1. Total Revenue (Net Sales)

How much we earned in total before cost deductions

```
print("Total Revenue: ", total_revenue)
```

```
Total Revenue: 9747747.933999998
```

### 4.2. Total Number of Paying Customers

Size of our customer base

```
unique_customers = df_clean['CustomerID'].nunique()
print(f"Total Unique Customers: {unique_customers}")
```

```
Total Unique Customers: 4338
```

### 4.3. Average Order Value (AOV)

Average amount spent per invoice

```
order_value = df_clean.groupby('InvoiceNo')['Revenue'].sum()
average_order_value = order_value.mean()
print(f"Average Order Value: {average_order_value:.2f}")
```

```
Average Order Value: 534.40
```

### 4.4. Monthly Revenue Trend

Are we growing or declining month-by-month?

```
df_clean['InvoiceDate'] = pd.to_datetime(df_clean['InvoiceDate']) #
ensure datetime format
df_date_indexed = df_clean.set_index('InvoiceDate') # set InvoiceDate
as index
monthly_revenue = df_date_indexed['Revenue'].resample('M').sum()
```

```
print("Monthly Revenue Breakdown:")
print(monthly_revenue)
```

```
# Plot
monthly_revenue.plot(kind='line', title='Monthly Revenue Trend',
ylabel='Revenue (£)', figsize=(10, 5))
```

```
Monthly Revenue Breakdown:
InvoiceDate
2010-12-31      823746.140
```

```


2011-01-31    691364.560
2011-02-28    523631.890
2011-03-31    717639.360
2011-04-30    537808.621
2011-05-31    770536.020
2011-06-30    761739.900
2011-07-31    719221.191
2011-08-31    759138.380
2011-09-30    1058590.172
2011-10-31    1154979.300
2011-11-30    1509496.330
2011-12-31     638792.680
Freq: M, Name: Revenue, dtype: float64

```

```

<Axes: title={'center': 'Monthly Revenue Trend'},
xlabel='InvoiceDate', ylabel='Revenue (£)'>

```



## 4.5. Top 10 Countries by Revenue

Where are our best customers located?

```

# Step 1: Calculate total revenue
total_revenue = df_clean['Revenue'].sum()

# Step 2: Group revenue by country
country_revenue = df_clean.groupby('Country')
['Revenue'].sum().sort_values(ascending=False)

```

```

# Step 3: Get top 10 + combine the rest into 'Others'
top_10 = country_revenue.head(10)
others = country_revenue[10:].sum()
top_10_with_others = top_10.copy()
top_10_with_others['Others'] = others

# Step 4: Calculate percentage share
country_percent = (top_10_with_others / total_revenue) * 100

# Step 5: Print country-wise revenue + percentage
print("Top 10 Countries + Others - Revenue & Percentage Share:\n")
for country, revenue in top_10_with_others.items():
    percent = country_percent[country]
    print(f"{country:<20} £{revenue:,.2f}    ({percent:.2f}%)")

# Step 6: Plot bar chart
plt.figure(figsize=(10, 6))
top_10_with_others.plot(kind='bar', color='darkcyan')
plt.title('Top 10 Countries + Others by Revenue')
plt.ylabel('Revenue (£)')
plt.xlabel('Country')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Step 7: Pie chart for "Others" countries only
# Get all countries after top 10
other_countries = country_revenue[10:]


# Optional: pick top 10 within "Others" for clarity, rest into
'Remaining'
top_from_others = other_countries.head(10).copy()
top_from_others['Remaining'] = other_countries[10:].sum()

# Step 8: Plot pie chart for Others only
plt.figure(figsize=(8, 8))
top_from_others.plot.pie(
    autopct='%1.1f%%',
    startangle=140,
    shadow=True,
    wedgeprops={'width': 0.5},
    cmap='tab20'
)
plt.title("Revenue Breakdown within 'Other' Countries")
plt.ylabel('')
plt.tight_layout()
plt.show()


```

Top 10 Countries + Others - Revenue & Percentage Share:

United Kingdom	£9,025,222.08	(84.61%)
Netherlands	£285,446.34	(2.68%)
EIRE	£283,453.96	(2.66%)
Germany	£228,867.14	(2.15%)
France	£209,715.11	(1.97%)
Australia	£138,521.31	(1.30%)
Spain	£61,577.11	(0.58%)
Switzerland	£57,089.90	(0.54%)
Belgium	£41,196.34	(0.39%)
Sweden	£38,378.33	(0.36%)
Others	£297,216.92	(2.79%)



Revenue Breakdown within 'Other' Countries



## 4.6. Top 10 Products by Revenue

Which products are performing best?

```
top_products = df_clean.groupby('Description')
['Revenue'].sum().sort_values(ascending=False)
print("Top 10 Products by Revenue:")
print(top_products.head(10))

(top_products.head(10)).plot(kind='barh', title='Top Selling Products
by Revenue', xlabel='Revenue (£)', figsize=(10, 6))
```

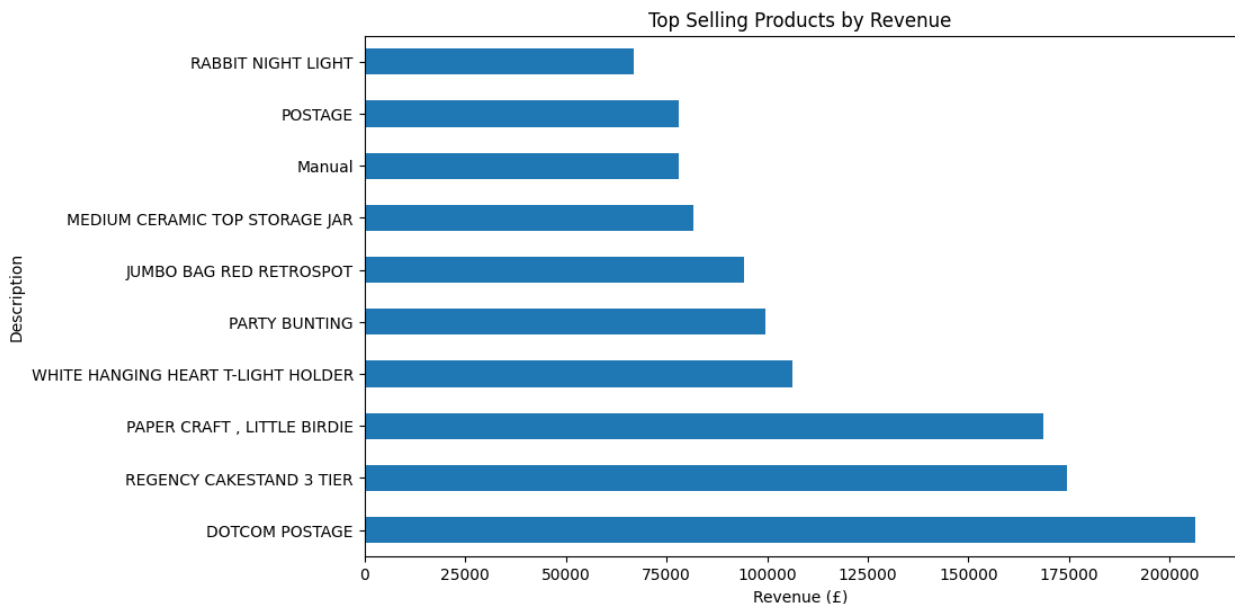


Top 10 Products by Revenue:

Description	
DOTCOM POSTAGE	206248.77
REGENCY CAKESTAND 3 TIER	174484.74
PAPER CRAFT , LITTLE BIRDIE	168469.60
WHITE HANGING HEART T-LIGHT HOLDER	106292.77
PARTY BUNTING	99504.33
JUMBO BAG RED RETROSPOT	94340.05
MEDIUM CERAMIC TOP STORAGE JAR	81700.92
Manual	78112.82
POSTAGE	78101.88
RABBIT NIGHT LIGHT	66964.99

Name: Revenue, dtype: float64

```
<Axes: title={'center': 'Top Selling Products by Revenue'},  
xlabel='Revenue (£)', ylabel='Description'>
```



## 4.7. Repeat vs One-Time Customers

How many customers returned to buy again?

```
customer_freq = df_clean['CustomerID'].value_counts()  
repeat_customers = (customer_freq > 1).sum()  
new_customers = (customer_freq == 1).sum()  
  
print(f"Repeat Customers: {repeat_customers}")  
print(f"One-Time Customers: {new_customers}")  
  
# Pie chart  
plt.figure(figsize=(6, 6))
```

```
plt.pie([repeat_customers, new_customers],
        labels=['Repeat Customers', 'One-Time Customers'],
        autopct='%1.1f%%', startangle=140, colors=['#66b3ff',
        '#ff9999'])
plt.title("Customer Loyalty Distribution")
plt.axis('equal')
plt.show()
```

Repeat Customers: 4267  
One-Time Customers: 71



## 4.8. Top 10 Customers by Lifetime Revenue

Who are our most valuable customers?

```
top_customers = df_clean.groupby('CustomerID')
['Revenue'].sum().sort_values(ascending=False).head(10)
print("Top 10 Customers by Revenue:")
print(top_customers)
```

```
top_customers.plot(kind='bar', title='Top Customers by Total Spend',
ylabel='Revenue (£)', figsize=(10, 5))
```

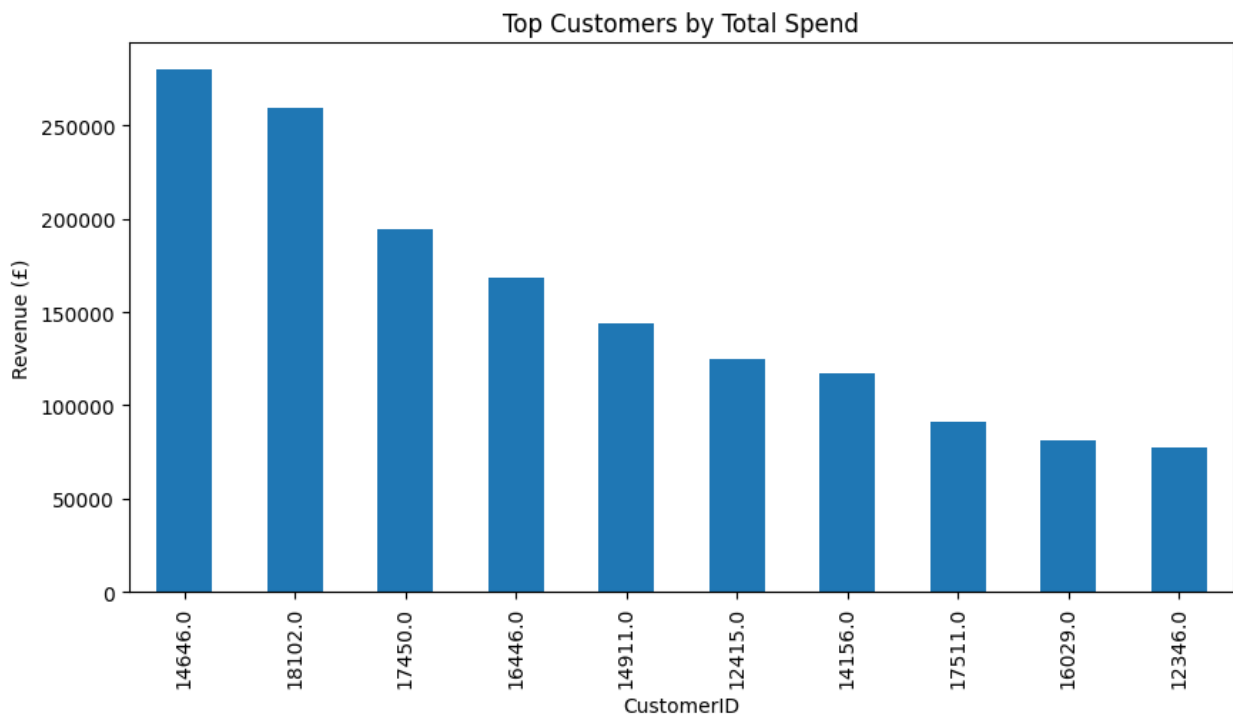
Top 10 Customers by Revenue:

CustomerID

14646.0	280206.02
18102.0	259657.30
17450.0	194550.79
16446.0	168472.50
14911.0	143825.06
12415.0	124914.53
14156.0	117379.63
17511.0	91062.38
16029.0	81024.84
12346.0	77183.60

Name: Revenue, dtype: float64

```
<Axes: title={'center': 'Top Customers by Total Spend'},
xlabel='CustomerID', ylabel='Revenue (£)'\>
```



## 4.9. Customer Lifetime Value

*# Total revenue per customer*

```
customer_revenue = df_clean.groupby('CustomerID')['Revenue'].sum()
```

*# Number of purchases per customer*

```
customer_orders = df_clean.groupby('CustomerID')
['InvoiceNo'].nunique()
```

```

# CLV = total revenue / total orders
clv = (customer_revenue /
customer_orders).sort_values(ascending=False)

# Display top 10 customers by CLV
print("Top 10 Customers by Lifetime Value (Simplified):")
print(clv.head(10))

# Plot top 10
clv.head(10).plot(kind='bar', title='Top 10 Customer Lifetime Value
(CLV)', ylabel='Avg Revenue per Order (£)', figsize=(10, 5))

```

Top 10 Customers by Lifetime Value (Simplified):


CustomerID	
16446.0	84236.250000
12346.0	77183.600000
15749.0	14844.766667
15098.0	13305.500000
12357.0	6207.670000
12415.0	5948.310952
12590.0	4932.130000
12688.0	4873.810000
12752.0	4366.780000
18102.0	4327.621667

dtype: float64

```

<Axes: title={'center': 'Top 10 Customer Lifetime Value (CLV)'},
xlabel='CustomerID', ylabel='Avg Revenue per Order (£)'>

```



THE END